

Then we will use the function `random.shuffle` to shuffle the list `[0, 1, ..., n-1]`, where `n` is the length of the data set. The `shuffle` function randomly rearranges the values of its argument in place, much like shuffling a deck of cards. We will call the list of indices `ind`, both before and after shuffling.

```
import random
n=len(data)
ind=list(range(n))
random.shuffle(ind)
```

At this point, the list `ind` still contains each of the integers `[0, 1, ..., n-1]`, but they have been rearranged in a random order. Since there are `n` total indices, then 75% of the indices (rounded to the nearest integer) is given by `ntrain=round(0.75*n)`. The idea is to extract the data pairs with indices given by the first `ntrain` numbers in the shuffled `ind` as the training set, and use the remaining data as the test set. Python lets us do this as follows:

```
ntrain=round(.75*n)
xtrain=xvals[ind[:ntrain]]
ytrain=yvals[ind[:ntrain]]
xtest=xvals[ind[ntrain:]]
ytest=yvals[ind[ntrain:]]
```

In this code, `ind[:ntrain]` gives the first 75% of the indices in the array `ntrain`, while `ind[ntrain:]` gives the remaining 25%. There is no-overlap, and all indices are included. By referring to `xvals[list_of_indices]` Python picks out those values in `xvals`, and only those values in `xvals`, that are expressed in `list_of_indices`. Here we have replaced `list_of_indices` with `ind[:ntrain]` to get a training set, and `ind[ntrain:]` to get a test set.

The two code blocks above can be replaced by the function in the `sklearn` package `model_selection.train_test_split`.¹⁵

```
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=\
    train_test_split(xvals.reshape(-1,1),yvals)
```

First, we'll use our old friend `scipy.stats.linregress` on the training set:

```
from scipy.stats import linregress as LR
b,a,r,p,stderr=LR(xtrain,ytrain)
print(a,b,r)
```

```
-0.2168130767260834 0.9246615087756963 0.9653789783330483
```

This tells us that the regression on the training set gives (rounded to 2 decimals)

$$ERP = 0.92 - 0.22 \times PRP \quad (1.67)$$

¹⁵For details see https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html. The `reshape` method ensure the that x-array is a column vector.